

```
In [1]: # Importing necessary libraries
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
```

```
In [10]: df = pd.read_csv(r"C:\Users\hp\Documents\Wine store data sheet.csv")
df.head()
```

```
Out[10]:
```

	winery	wine	year	rating	num_reviews	country	region	price	type	body	acidity
0	Teso La Monja	Tinto	2013	4.9	58	Espana	Toro	995.00	Toro Red	5.0	3.0
1	Artadi	Vina El Pison	2018	4.9	31	Espana	Vino de Espana	313.50	Tempranillo	4.0	2.0
2	Vega Sicilia	Unico	2009	4.8	1793	Espana	Ribera del Duero	324.95	Ribera Del Duero Red	5.0	3.0
3	Vega Sicilia	Unico	1999	4.8	1705	Espana	Ribera del Duero	692.96	Ribera Del Duero Red	5.0	3.0
4	Vega Sicilia	Unico	1996	4.8	1309	Espana	Ribera del Duero	778.06	Ribera Del Duero Red	5.0	3.0

```
In [11]: #Data Cleaning and validation
df.isnull().sum()
```

```
Out[11]: winery      0
wine      0
year      2
rating    0
num_reviews  0
country   0
region    0
price     0
type     545
body     1169
acidity   1169
dtype: int64
```

```
In [12]: #Dropping the null values on our dataset
df = df.dropna()
```

```
In [13]: df.shape
```

```
Out[13]: (6329, 11)
```

```
In [17]: df.describe()
```

Out[17]:

	rating	num_reviews	price	body	acidity
count	6329.000000	6329.000000	6329.000000	6329.000000	6329.000000
mean	4.259425	442.292463	65.659082	4.158319	2.946753
std	0.124306	718.597235	162.599997	0.583345	0.247955
min	4.200000	25.000000	4.990000	2.000000	1.000000
25%	4.200000	388.000000	19.980000	4.000000	3.000000
50%	4.200000	402.000000	29.150000	4.000000	3.000000
75%	4.200000	415.000000	60.950000	5.000000	3.000000
max	4.900000	32624.000000	3119.080000	5.000000	3.000000

In [16]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6329 entries, 0 to 7499
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   winery           6329 non-null   object
1   wine             6329 non-null   object
2   year             6329 non-null   object
3   rating           6329 non-null   float64
4   num_reviews      6329 non-null   int64
5   country          6329 non-null   object
6   region           6329 non-null   object
7   price            6329 non-null   float64
8   type             6329 non-null   object
9   body             6329 non-null   float64
10  acidity           6329 non-null   float64
dtypes: float64(4), int64(1), object(6)
memory usage: 593.3+ KB
```

```
In [19]: #I'M REMOVING THE YEAR BECAUSE IT CONTAIN ALOT OF NV MAKING IT DIFFICUT TO CONVERT
df['year'] = df['year'].replace('N.V.', np.NaN)
df = df.dropna()
df['year'] = df['year'].astype(np.int64)
```


```
In [20]: #The column country only has one value so it would not be helpful at all for our model
df = df.drop(columns=['country'])
df.head()
```

Out[20]:

	winery	wine	year	rating	num_reviews	region	price	type	body	acidity
0	Teso La Monja	Tinto	2013	4.9	58	Toro	995.00	Toro Red	5.0	3.0
1	Artadi	Vina El Pison	2018	4.9	31	Vino de Espana	313.50	Tempranillo	4.0	2.0
2	Vega Sicilia	Unico	2009	4.8	1793	Ribera del Duero	324.95	Ribera Del Duero Red	5.0	3.0
3	Vega Sicilia	Unico	1999	4.8	1705	Ribera del Duero	692.96	Ribera Del Duero Red	5.0	3.0
4	Vega Sicilia	Unico	1996	4.8	1309	Ribera del Duero	778.06	Ribera Del Duero Red	5.0	3.0


In []: *#OUR DATA IS CLEAN AND READY FOR Exploratory Data Analysis*In [21]: `sns.heatmap(df.corr(), annot=True, cmap='Blues')`

Out[21]: <AxesSubplot:>



In [22]: *#WE WANT TO CHECK IF THE WINE AFFECTS THE WINE PRICE*

```
fig, ax = plt.subplots(ncols=1, figsize=(18,7))
sns.boxplot(y='price', x='type', data=df, ax=ax)
plt.xticks(rotation=90)
plt.show()
```



```
In [23]: print('Categorical columns: ')
for col in df.columns:
    if df[col].dtype == 'object':
        print(str(col))
        label = LabelEncoder()
        label = label.fit(df[col])
        df[col] = label.transform(df[col].astype(str))
```

Categorical columns:
winery
wine
region
type

```
In [24]: df = (df-df.mean())/df.std()
df.head()
```

```
Out[24]:
```

	winery	wine	year	rating	num_reviews	region	price	type	body
0	1.281015	1.291723	0.006271	5.114070	-0.631438	1.143410	5.604341	2.208328	1.407862
1	-1.766666	1.546400	0.706038	5.114070	-0.676061	1.355532	1.486888	1.932690	-0.275693
2	1.493195	1.374770	-0.553542	4.314853	2.235989	0.294921	1.556066	0.003224	1.407862
3	1.493195	1.374770	-1.953074	4.314853	2.090552	0.294921	3.779491	0.003224	1.407862
4	1.493195	1.374770	-2.372934	4.314853	1.436084	0.294921	4.293644	0.003224	1.407862

```
In [27]: #Splitting the data using train test split
X = df.drop(columns=['price'])
y = df['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2)
```

```
In [28]: #import models for regression
from sklearn.linear_model import LinearRegression, Lasso, Ridge, BayesianRidge
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import LinearSVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [29]: models = {}
def train_validate_predict(regressor, x_train, y_train, x_test, y_test, index):
    model = regressor
    model.fit(x_train, y_train)

    y_pred = model.predict(x_test)

    r2 = r2_score(y_test, y_pred)
    models[index] = r2
```

```
In [42]: model_list = [LinearRegression, Lasso, Ridge, BayesianRidge, DecisionTreeRegressor,
    RandomForestRegressor]
model_names = ['Linear Regression', 'Lasso', 'Ridge', 'Bayesian Ridge', 'Decision Tree Regressor',
    'KNeighbors Regressor', 'Random Forest Regressor']
```


```
In [43]: models
```

```
Out[43]: {'Linear Regression': 0.4181414658706233,
'Lasso': -0.0015752423477575217,
'Ridge': 0.4181350218488735,
'Bayesian Ridge': 0.41800899693022064,
'Decision Tree Regressor': 0.2772944875167366,
'Linear SVR': 0.1950185249605486,
'KNeighbors Regressor': 0.7190891654351761,
'Random Forest Regressor': 0.6417788308955521}
```

```
In [33]: #Evaluating
model = KNeighborsRegressor()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
preds = pd.DataFrame({'y_pred': y_pred, 'y_test': y_test})
preds = preds.sort_values(by='y_test')
preds = preds.reset_index()
```

```
In [34]: #VISUALIZATION
plt.figure(figsize=(15, 5))
plt.plot(preds['y_pred'], label='pred')
plt.plot(preds['y_test'], label='actual')
plt.legend()
plt.show()
```



```
In [ ]: # columns has a very little to no relationship toward to the wines prices.

#The model did alright at predicting low prices wines but did terrible at high price
#I think what caused this from happening according to our EDA earlier, that in our
#theres way more data on low prices wines but theres a little data from the high price
```